

Network Programming

Nathan Warner



Northern Illinois
University

Computer Science
Northern Illinois University
United States

Contents

1	A brief overview	4
2	Network Terminology	4
2.1	Nodes, links, and paths	4
2.2	Networking protocol	5
3	The OSI (Opens systems interconnection) Model	6
3.1	Main idea	7
3.2	Physical layer: Wired media	8
3.3	Physical layer: Wireless media	8
3.4	Data link layer: Functionalities	9
3.4.1	Example: Ethernet frame	9
3.4.2	Example: Data link flow	9
3.5	Network layer (Internet protocol layer)	9
3.5.1	IPv4 Address	10
3.5.2	IPv6 Address	10
3.5.3	IP Packet	10
3.6	IP Layer: Routing and Forwarding	11
3.7	Transport layer	11
3.8	Ports	11
3.8.1	Transport layer programming	11
3.8.2	Transport layer protocols	12
4	Domain Name Service (DNS)	13
4.1	Domain Names	13
4.2	C library function: getaddrinfo	13
4.2.1	Signature	13

4.2.2	Address info structure	13
4.2.3	Socket address info structure	14
4.2.4	Example: Get host by name	15
4.3	gai_strerror	16
4.3.1	Signature	16
4.3.2	ecodes	16
4.3.3	Return value	16

5 User Datagram Protocol (UDP) 17

5.1	UDP Packet Format	17
5.2	UDP Programming	17
5.2.1	Socket system calls	17
5.2.2	UDP communications pattern	18
5.3	System call: Socket	18
5.3.1	Signature	18
5.4	System call: Bind	18
5.4.1	Signature	18
5.4.2	Structure sockaddr: 16 bytes	19
5.4.3	sockaddr_in members we care about	19
5.5	Byte order	19
5.6	Helper functions: htons, htonl, ntohs, ntohl	20
5.6.1	Signatures	20
5.7	Helper function: inet_addr	20
5.7.1	Signature	21
5.7.2	Return values	21
5.8	System call: recvfrom	21
5.8.1	Signature	21
5.9	System call: sendto	21
5.9.1	Signature	22
5.10	System call: close	22
5.10.1	signature	22
5.11	UDP Programming example: simple server - echo	23
5.12	UDP programming example: Simple client - echo	25

6	Transmission Control Protocol (TCP)	27
6.1	TCP / IP Protocol packet	27
6.2	TCP Communication	27
6.3	Socket System Calls	28
6.4	TCP Communications Pattern	28
6.5	System call: Socket	28
6.5.1	Signature	28
6.6	Client System call: connect	29
6.6.1	Signature	29
6.7	TCP Client illustration	29
6.8	Server system call: bind	32
6.8.1	Signature	32
6.9	Server system call: Listen	32
6.9.1	Signature	32
6.10	Server system call: Accept	32
6.10.1	Signature	32
6.11	TCP Server illustration	33
7	TCP Server & Shell Job Control	36
7.1	Improve TCP Client	36
7.1.1	Accept FQDN as server address	36
7.1.2	Process complete server response	36
7.1.3	Full example	37
7.2	Improve TCP server	39
7.2.1	TCP Server Fork	39
7.2.2	TCP Server Fork basic logic	39
7.2.3	TCP Server Fork example	40
7.3	Improved TCP Server and Client example: List directory	43
7.3.1	Program	44

A brief overview

In the context of network programming, "networking" refers to the practice and techniques involved in designing, implementing, and managing communication between computers and devices over a network. This can include a wide array of tasks and principles, including but not limited to:

- **Data Communication:** The fundamental aspect of networking, involving the exchange of data between two or more devices over a network. This can be achieved through various communication protocols and standards.
- **Protocols and Standards:** Networking relies on a set of rules and conventions (protocols) for communication between network devices. These protocols define how data is formatted, transmitted, and received. Examples include TCP/IP (Transmission Control Protocol/Internet Protocol), HTTP (HyperText Transfer Protocol), and FTP (File Transfer Protocol).
- **Network Architecture:** The design and layout of a network, including its components (e.g., routers, switches, gateways) and topology (e.g., star, mesh, ring). Network architecture decisions impact the network's performance, scalability, and security.
- **Socket Programming:** A means of connecting two nodes on a network to communicate with each other. One node listens on a particular port at an IP, while another node connects to it. Socket programming is used to facilitate communication between applications running on different computing devices.
- **APIs for Network Communication:** Programming interfaces such as Winsock for Windows, POSIX sockets for Unix/Linux, and various cross-platform networking libraries (e.g., Boost.Asio) that allow developers to implement networking functionalities.
- **Network Services Development:** Creating software that provides specific functionalities over a network, such as web servers, email servers, and file sharing systems.
- **Network Security:** Ensuring the confidentiality, integrity, and availability of data in the network. This includes implementing secure protocols (like HTTPS), encryption, firewalls, and intrusion detection systems.
- **Network Management:** Monitoring and maintaining network operations. This involves performance analysis, troubleshooting network problems, and ensuring that network resources are allocated efficiently.

Network Terminology

2.1 Nodes, links, and paths

node refers to any device that can send, receive, or forward information over a communications channel. Nodes can be computers, mobile devices, routers, switches, and other devices capable of processing or storing data.

A link, on the other hand, is the physical or logical connection between two or more nodes, enabling them to communicate. Links can be wired connections like Ethernet cables, optical fibers, or wireless connections such as Wi-Fi or Bluetooth.

Together, nodes and links form the basic components of a network, allowing for the transmission of data across diverse and complex systems.

A path is a sequence of nodes and links

In other words...

- **Node:** Host or intermediary
- **Link:** Point-to-point or broadcast to many other nodes at the same time
- **Link medium:** wired or wireless
- **Path:** Routed or switched (Elaborated in later section)

2.2 Networking protocol

Networking protocols are standardized sets of rules that determine how data is transmitted and received across a network. These protocols specify the formats for data packets, the procedures for signaling, error handling, and data encryption to ensure successful communication between devices.

More broadly, information is exchanged between nodes via **messages**, each message has an exact meaning intended to provoke a defined response of the receiver

Note:-

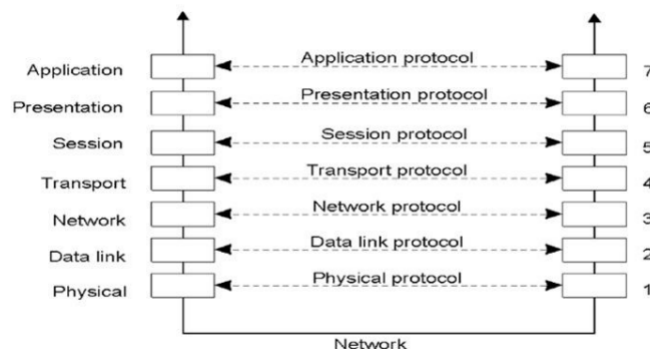
Messages used **well-defined format**

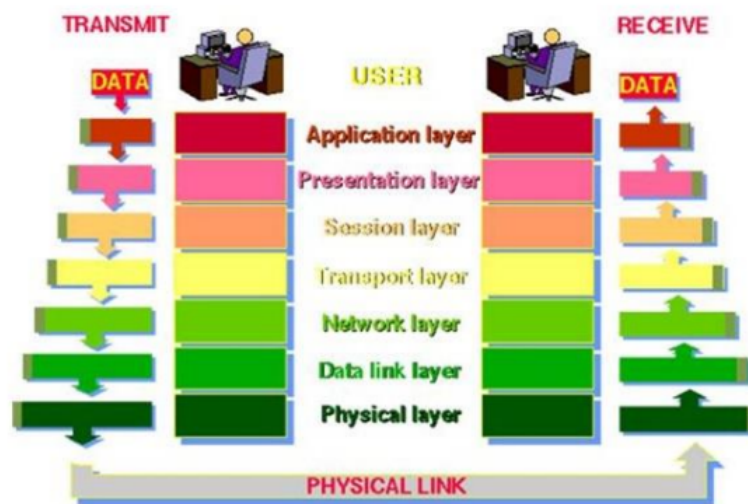
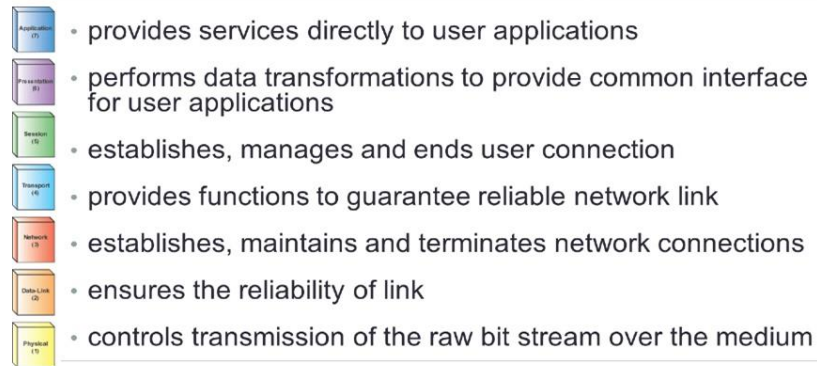
The OSI (Opens systems interconnection) Model

The OSI (Open Systems Interconnection) model is a conceptual framework used to understand and standardize the functions of a telecommunications or computing system without regard to its underlying internal structure and technology. Developed by the International Organization for Standardization (ISO), the OSI model divides the process of communication between two end-points in a network into seven layers. Each layer serves a specific function and communicates with the layers directly above and below it. From top to bottom, the layers are **APS TNDP**:

- **Application Layer (Layer 7):** The closest to the end user, this layer interacts with software applications that implement a communicating component. It provides protocols that allow software to send and receive information and present meaningful data to the user (e.g., HTTP for web browsing, SMTP for email).
- **Presentation Layer (Layer 6):** Translates data between the application layer and the network format. It ensures data is in a usable format and can encrypt or compress data if necessary.
- **Session Layer (Layer 5):** Manages sessions between applications, establishing, managing, and terminating connections between local and remote applications.
- **Transport Layer (Layer 4):** Responsible for data transfer between end systems and provides reliable data transfer services to the upper layers. This includes breaking down messages into smaller units if needed, and ensuring error-free data transfer (e.g., TCP and UDP).
- **Network Layer (Layer 3):** Manages device addressing, identifies the best paths for data transmission, and routes data packets between devices that are not locally attached. Routers operate at this layer.
- **Data Link Layer (Layer 2):** Provides data transfer across the physical link established by the physical layer. It deals with MAC addresses, error detection and correction, and frames data packets.
- **Physical Layer (Layer 1):** Concerns the physical equipment involved in data transfer, such as cables, switches, and the electrical signals that traverse these media.

In other words, these seven layers help to describe communications in a network





3.1 Main idea

The complexities of communication is organized into successive layers of protocols

- **Lower-level layers:** Specific to medium
- **Higher-level layers:** Specific to application

3.2 Physical layer: Wired media

- **Ethernet**
 - 10BASE-T, 100BASE-T, 1000BASE-T
 - 10GbE, 40GbE, 100GbE
- **Business/backbone**
 - DS1(T1): 1.54Mbps to DS5: 400Mbps
 - OC-1: 50Mbps to OC-768: 40Gbs
- **Last mile:**
 - Modem
 - DSL (Digital subscriber lines)
 - Cable: DOCSIS
 - FiOS (Fiber optic service)

3.3 Physical layer: Wireless media

- **Cellphone Data:**
 - EDGE, GPRS, HSPA+
 - 4G LTE up to 100Mbps
 - 5G over 100Mbps
- **Satellite**
 - Wildblue: 12Mbps
 - Hughesnet: 15Mbps
 - Starlink: 200Mbps
- **WiFi: 802.11**
 - Up to 150Mbps & MIMO
 - New: "ac" up to 1Gbs
- **WiMax: 802.16**
 - up to 40Mbps
- **WPAN**
 - Bluetooth up to 2Mbps
 - NFC up to 423Kbs
 - ZigBee up to 256Kbs

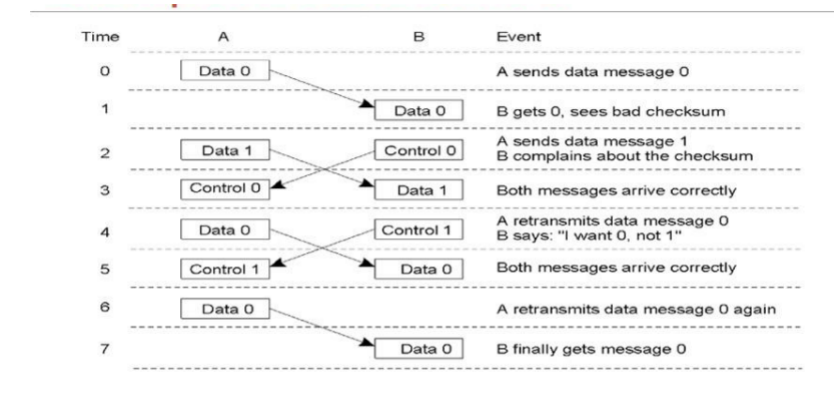
3.4 Data link layer: Functionalities

- **Medium access control:** Arbitrate who transmits
- **Addressing:** address of receiver, address of sender
- **Framing:** Delimited unit of transmission for data & control
- **Error control and reliability**
- **Flow control**

3.4.1 Example: Ethernet frame

Preamble	Destination MAC address	Source MAC address	Type/Length	User Data	Frame Check Sequence (FCS)
8	6	6	2	46 - 1500	4

3.4.2 Example: Data link flow



3.5 Network layer (Internet protocol layer)

Provides host to host transmission service, where hosts are not necessarily adjacent

Layer provides services

- **Addressing**
 - Hosts have global addresses: IPv4, IPv6
 - Uses data link layer protocol to translate address
- **Routing and forwarding:** Find path from host to host

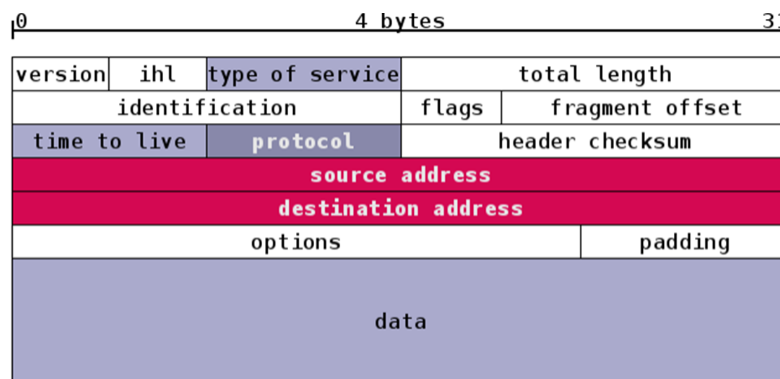
3.5.1 IPv4 Address

•IP address	•127.0.0.1
•32bit unique identifier, written as quad	•131.156.145.90
•network	
•first n bits of IP number, written as "/n"	•131.156.0.0/16
•8 - class A, 16 - class B, 24 - class C	•131.156.145.0/24
•more than 24 - class D	
•netmask	
•32 bit number with first n bits all 1, rest 0	•255.255.255.0
•broadcast	
•network number (first n bits), rest all 1	•131.156.145.255
•gateway IP	
•name server IP	•131.156.145.1
	•131.156.145.2

3.5.2 IPv6 Address

-
- IP address: 128-bit unique identifier
 - 8 groups of 16-bit values,
each group in 4 hex digits, separated by ":"
 - ex.: 2001:0db8:0000:0000:0000:ff00:0042:8329
 - can be abbreviated:
 - remove leading zeroes: 42 instead of 0042
 - omit consecutive sections of zeroes: 2001:db8::ff00:42:8329

3.5.3 IP Packet



3.6 IP Layer: Routing and Forwarding

Done by hosts on path from sending to receiver

- **Forwarding:** Host has 2 network interfaces, transfers packet from incoming to outgoing interface
- **Routing:**
 - Finds path from sender to receiver
 - **Simple routing:** know receiver or send to gateway
 - **Advanced routing:** determine which gateway to send to (typically with multiple outgoing network interfaces)

3.7 Transport layer

Provides end-to-end communication services for applications

Byte format as abstraction on underlying system format. Raises reliability

Also enables multiplexing, which provides multiple endpoints on a single node: **port**.

Refines connection address via port number

3.8 Ports

-
- 0 to 1023: well-known ports
 - 20 & 21: File Transfer Protocol (FTP)
 - 22: Secure Shell (SSH)
 - 23: Telnet remote login service
 - 25: Simple Mail Transfer Protocol (SMTP)
 - 53: Domain Name System (DNS) service
 - 80: Hypertext Transfer Protocol (HTTP) used in the World Wide Web
 - 110: Post Office Protocol (POP3)
 - 119: Network News Transfer Protocol (NNTP)
 - 143: Internet Message Access Protocol (IMAP)
 - 161: Simple Network Management Protocol (SNMP)
 - 443: HTTP Secure (HTTPS)
 - 1024 to 49151: IANA registered ports
 - 49152 to 65535: dynamic or private port

3.8.1 Transport layer programming

- **Common abstraction:** Socket
- Socket is end-point of communication link
 - Identified as Ip address + port number
- operates as client and server

3.8.2 Transport layer protocols

-
- TCP: transmission control protocol
 - connection oriented, guaranteed delivery
 - stream oriented: basis for: http, ftp, smtp, ssh
 - UDP: user datagram protocol
 - best effort
 - datagram oriented: basis for: dns, rtp
 - DCCP: datagram congestion control protocol
 - SCTP: stream control transmission protocol

Domain Name Service (DNS)

4.1 Domain Names

- hierarchical distributed naming system
- Uses **FQDN**: Fully qualified domain name
- **DNS**: Domain name service
 - Resolves query for FQDN into ip address

4.2 C library function: getaddrinfo

4.2.1 Signature

```
1  int getaddrinfo(const char* node, const char* service, const  
    ↪ struct addrinfo* hints, struct addrinfo** res)
```

- Translates FQDN **node** into IP address
- **res** is pointer to a list of address info structures

Note:-

service and **hints** can be NULL

4.2.2 Address info structure

```
1  struct addrinfo {  
2      int             ai_flags;  
3      int             ai_family;  
4      int             ai_socktype;  
5      int             ai_protocol;  
6      size_t          ai_addrlen;  
7      struct sockaddr *ai_addr; // Socket address  
8      char            ai_cannonname;  
9      struct addrinfo *ai_next;  
10 };
```

4.2.3 Socket address info structure

```
1  struct sockaddr_in {
2      short                sin_family    // e.g. AF_INET
3      unsigned_short       sin_port     // Port
4      struct in_addr       sin_addr     // IP Address
5      char                 sin_zero[8]  // Padding
6  };
7
8  struct in_addr {
9      unsigned long        s_addr;
10 };
```

Note:-

sin_addr can be printed via **inet_ntoa** fn

4.2.4 Example: Get host by name

```
1  /*
2  * getHostName.cxx
3  *
4  *      do DNS lookup
5  *
6  */
7  #include <sys/types.h>
8  #include <sys/socket.h>
9  #include <netinet/in.h>
10 #include <arpa/inet.h>
11 #include <netdb.h>
12 #include <cstdio>
13 #include <cstdlib>
14 #include <iostream>
15 using namespace std;
16
17 int main(int argc, char*argv[]) {
18     struct addrinfo *res;
19     int error;
20     const char *hostname = "faculty.cs.niu.edu";
21
22     if (argc > 1) {
23         hostname = argv[1];
24     }
25
26     error = getaddrinfo(hostname, NULL, NULL, &res);
27     if (error) {
28         cerr << hostname << ": " << gai_strerror(error) << endl;
29         exit(EXIT_FAILURE);
30     }
31
32     // convert generic sockaddr to Internet sockaddr_in
33     struct sockaddr_in *addr = (struct sockaddr_in *)
↪ res->ai_addr;
34     // convert network representation into printable presentation
35     cout << hostname << " is: " << inet_ntoa(addr->sin_addr) <<
↪ endl;
36 }
```


4.3 gai_strerror

Concept 1: The `gai_strerror()` function shall return a text string describing an error value for the `getaddrinfo()` and `getnameinfo()` functions listed in the `<netdb.h>` header.

4.3.1 Signature

```
1  const char* gai_strerror(int ecode);
```

4.3.2 ecodes

When the `ecode` argument is one of the following values listed in the `<netdb.h>` header:

[EAI_AGAIN]
[EAI_BADFLAGS]
[EAI_FAIL]
[EAI_FAMILY]
[EAI_MEMORY]
[EAI_NONAME]
[EAI_OVERFLOW]
[EAI_SERVICE]
[EAI_SOCKTYPE]
[EAI_SYSTEM]

the function return value shall point to a string describing the error. If the argument is not one of those values, the function shall return a pointer to a string whose contents indicate an unknown error.

4.3.3 Return value

Upon successful completion, `gai_strerror()` shall return a pointer to an implementation-defined string.

User Datagram Protocol (UDP)

- Simple message-based connection-less protocol
 - Transmits information in one direction from source to destination without verifying the readiness or state of the receiver
- Uses **datagram** as message
- Stateless and fast

5.1 UDP Packet Format

bits	0 – 7	8 – 15	16 – 23	24 – 31
0	Source IP address			
32	Destination IP address			
64	Zeros	Protocol	UDP length	
96	Source Port		Destination Port	
128	Length		Checksum	
160+	Data			

5.2 UDP Programming

- **Common abstraction:** Socket
- Socket is end-point of communication link
 - Identified as IP address + port number
 - can receive data, can send data
- **Typical logic:** Server vs client
 - Server ready to receive datagram from any client
 - client sends datagram to specific server
 - server responds with datagram to client

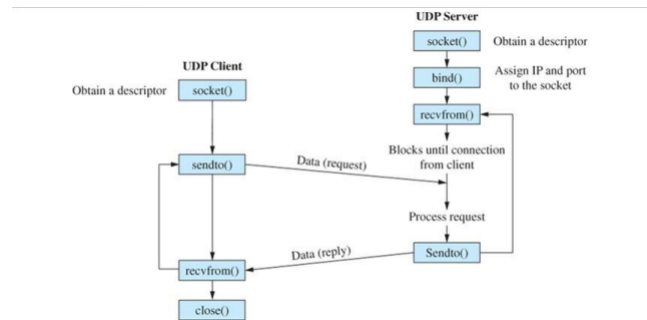
5.2.1 Socket system calls

server

System call	Meaning
socket	Create a new communication endpoint
bind	Attach a local address to a socket
recvfrom	Receive(read) some data over the connection
sendto	Send(write) some data over the connection
close	Release the connection

client

5.2.2 UDP communications pattern



5.3 System call: Socket

5.3.1 Signature

```
1 int socket(int domain, int type, int protocol)
```

- Creates a new socket, as end point to a communications link
- **Domain** is set to **AF_INET**
- **Type** is set to **SOCK_DGRAM** for datagrams
- **Protocol** is set to 0, i.e. default UDP
- Returns socket descriptor
 - Used in **bind**, **sendto**, **recvfrom**, **close** (System calls)

5.4 System call: Bind

5.4.1 Signature

```
1 int bind(int sockfd, const struct sockaddr* addr, socklen_t
  ↪  addrlen)
```

- **Assigns address to socket:** IP number and port
- **struct sockaddr** holds address information
 - Will accept **struct sockaddr_in** pointer
- **addrlen** specifies length of **addr** structure
- **Returns** 0 on success, -1 otherwise

5.4.2 Structure sockaddr: 16 bytes

```
1  struct sockaddr {
2      short      sa_family    // Address family
3      char       sa_data[14] // Address data
4  };
5
6  struct sockaddr_in {
7      short      sin_family    // address family
8      unsigned short sin_port    // port number: 2 bytes
9      struct in_addr sin_addr    // IP address: 4 bytes
10     char       sin_zero[8]
11 };
12
13 struct in_addr {
14     unsigned long s_addr;
15 };
```

5.4.3 sockaddr_in members we care about

- **sin_family**: Always `AF_INET` in ipv4 communications
- **sin_port**: We use `htons(n)` fn to convert a c++ integer to a 16bit unsigned integer.
e.g. `htons(4444)`
- **sin_addr**: Either
 - `INADDR_ANY`: For server, to specify the server IP
 - `inet_addr("numeric ip (quad) as string")`: For client.

5.5 Byte order

In C programming, especially when dealing with network applications, it's crucial to manage byte order correctly due to differences in how data is stored in memory across different computer architectures. Byte order refers to the sequence in which bytes are arranged into larger data types (like integers) when stored in memory. The two most common byte orders are:

1. **Big-endian**: The most significant byte (MSB) is stored at the smallest address, the next significant byte in the next address, down to the least significant byte (LSB) at the highest address.
2. **Little-endian**: The LSB is at the smallest address, and the MSB is at the highest address.

5.6 Helper functions: htons, htonl, ntohs, ntohl

To ensure data consistency across different systems, certain helper functions are used in network programming to convert between host and network byte orders. These are part of the POSIX (Portable Operating System Interface) standards, available in `<arpa/inet.h>` in C.

- **htonl() (Host TO Network Long):** Converts a 32-bit integer from host byte order to network byte order.
- **htons() (Host TO Network Short):** Converts a 16-bit integer from host byte order to network byte order.
- **ntohl() (Network TO Host Long):** Converts a 32-bit integer from network byte order to host byte order.
- **ntohs() (Network TO Host Short):** Converts a 16-bit integer from network byte order to host byte order.

When sending data across a network, you would typically use these functions to convert multi-byte integers from the host's native byte order to network byte order before transmission. Similarly, upon receipt, you would convert the data from network byte order back to the host's native byte order. This handling ensures that applications developed on different platforms can communicate effectively over a network.

For example, if you're developing a network application that sends a 32-bit integer from a little-endian machine (like most Intel and AMD processors), you would use `htonl()` to convert the integer to big-endian format before sending it. On the receiving side, if the machine is big-endian, it would use `ntohl()` to convert the received integer back to its native byte order.

5.6.1 Signatures

```
1  uint16_t htons(uint16_t hostshort);
2  uint32_t htonl(uint32_t hostlong);
3  uint16_t ntohs(uint16_t netshort);
4  uint32_t ntohl(uint32_t netlong);
```

5.7 Helper function: inet_addr

The `inetb_addr` function is a commonly used helper function in C network programming, provided by the `<arpa/inet.h>` header file. Its primary purpose is to convert an IPv4 address in its standard text representation (a string in dotted-decimal format) into a numeric binary format, which is used in various network-related system calls and structures.

5.7.1 Signature

```
1  in_addr_t inet_addr(const char* cp)
```

Where **cp** is a constant character pointer to a null-terminated string representing the IPv4 address in dotted-decimal notation (e.g., "192.168.0.1").

5.7.2 Return values

- On success, **inet_addr** returns the IPv4 address as a `uint32_t` in **network byte order (big-endian)**.
- If the string in **cp** does not contain a valid IP address, the function returns **IN_ADDR_NONE** (usually defined as `(uint32_t) -1`).

5.8 System call: recvfrom

5.8.1 Signature

```
1  ssize_t recvfrom(int sockfd, void* buf, size_t len, int flags,  
    ↪ struct sockaddr* src_addr, socklen_t* addrlen)
```

- Receives a datagram **buf** of size **len** from socket **sockfd**
 - will wait until a datagram is available
 - **flags** specify wait behavior, e.g: 0 for default
- **src_addr** will hold address information of sender
 - **struct sockaddr** defines address structure
 - **addrlen** specifies length of **src_addr** structure
- Returns the number of bytes received, i.e size of datagram

5.9 System call: sendto

5.9.1 Signature

```
1 ssize_t sendto(int sockfd, const void* buf, size_t len, int
  ↪ flags, const struct sockaddr* dest_addr, socklen_t addrlen)
```

- Sends datagram **buf** of size **len** to socket **sockfd**
 - Will wait if there is no ready receiver
 - Flags specifies wait behavior, e.g: 0 for default
- **dest_addr** holds address information of receiver
 - **struct sockaddr** defines address structure
 - **addrlen** specifies length of **dest_addr** structure
- Returns the number of bytes sent, i.e size of datagram

5.10 System call: close

5.10.1 signature

```
1 int close(int fd)
```

- Closes socket specified by **fd** socket descriptor
- returns zero on success

5.11 UDP Programming example: simple server - echo

```
1  /*
2  * echoServer.cxx
3  *
4  * UDP echo server
5  *
6  *      loops/waits for message received from client
7  *      send message back to client
8  *
9  *      command line arguments:
10 *      argv[1] port number to receive from
11 *
12 */
13 #include <sys/socket.h>
14 #include <arpa/inet.h>
15 #include <unistd.h>
16
17 #include <cstdio>
18 #include <cstdlib>
19 #include <cstring>
20 #include <iostream>
21 using namespace std;
22
23 int main(int argc, char *argv[]) {
24
25     if (argc != 2) {
26         cerr << "USAGE: echoServer port\n";
27         exit(EXIT_FAILURE);
28     }
29
30     char buffer[256];
31     int received = 0;
32
33     int sock;
34     struct sockaddr_in server_address; // structure for address
    ↪ of server
35     struct sockaddr_in client_address; // structure for address
    ↪ of client
36     unsigned int addrlen = sizeof(client_address);
37
38     // Create the UDP socket
39     if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
40         perror("socket");
41         exit(EXIT_FAILURE);
42     }
43 }
```



```

1      // Construct the server sockaddr_in structure
2      memset(&server_address, 0, sizeof(server_address));    /*
↳ Clear struct */
3      server_address.sin_family = AF_INET;                    /*
↳ Internet/IP */
4      server_address.sin_addr.s_addr = INADDR_ANY;           /*
↳ Any IP address */
5      server_address.sin_port = htons(atoi(argv[1]));      /*
↳ server port */
6
7      // Bind the socket
8      if (bind(sock, (struct sockaddr *) &server_address,
↳ sizeof(server_address)) < 0) {
9          perror("bind");
10         exit(EXIT_FAILURE);
11     }
12
13     cout << "echoServer listening on port: " << argv[1] << endl;
14
15     // Run until cancelled
16     while (true) {
17         // Receive a message from the client
18         if ((received = recvfrom(sock, buffer, 256, 0, (struct
↳ sockaddr *) &client_address, &addrlen)) < 0) {
19             perror("recvfrom");
20             exit(EXIT_FAILURE);
21         }
22         cout << "Client (" << inet_ntoa(client_address.sin_addr)
↳ << ") sent " << received << " bytes: " << buffer << endl;
23         // Send the message back to client
24         if (sendto(sock, buffer, received, 0, (struct sockaddr
↳ *) &client_address, addrlen) < 0) {
25             perror("sendto");
26             exit(EXIT_FAILURE);
27         }
28     }
29
30     close(sock);
31     return 0;
32 }

```

5.12 UDP programming example: Simple client - echo

```
1  /*
2  * echoClient.cxx
3  *
4  * UDP echo client
5  *
6  *      sends message to echo server
7  *      waits for message received from server
8  *
9  *      command line arguments:
10 *      argv[1] IP number of server
11 *      argv[2] port number to send to
12 *      argv[3] message to send
13 *
14 */
15 #include <sys/socket.h>
16 #include <arpa/inet.h>
17 #include <unistd.h>
18
19 #include <cstdio>
20 #include <cstdlib>
21 #include <cstring>
22 #include <iostream>
23 using namespace std;
24
25 int main(int argc, char *argv[]) {
26
27     if (argc != 4) {
28         cerr << "USAGE: echoClient server_ip port
↪ message\n";
29         exit(EXIT_FAILURE);
30     }
31
32     char buffer[256];
33     int echolen, received = 0;
34
35     int sock;
36     struct sockaddr_in server_address; // structure for
↪ address of server
37     unsigned int addrlen = sizeof(server_address);
38
39     // Create the UDP socket
40     if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
41         perror("Failed to create socket");
42         exit(EXIT_FAILURE);
43     }
```

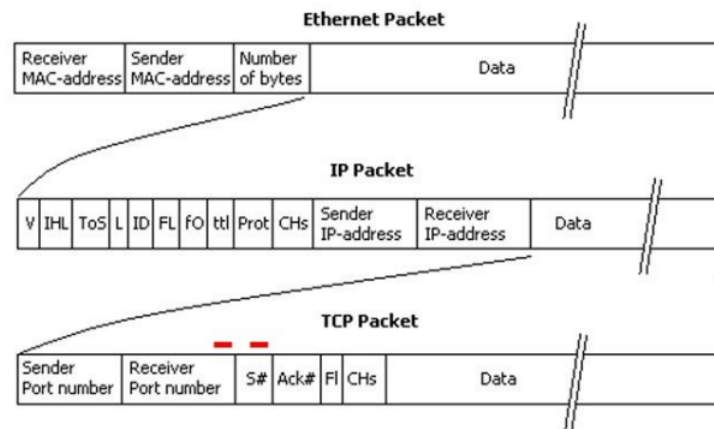
```

1      // Construct the server sockaddr_in structure
2      memset(&server_address, 0, sizeof(server_address));
↳ /* Clear struct */
3      server_address.sin_family = AF_INET;
↳ /* Internet/IP */
4      server_address.sin_addr.s_addr = inet_addr(argv[1]);
↳ /* IP address */
5      server_address.sin_port = htons(atoi(argv[2]));
↳ /* server port */
6
7      // Send the message to the server (don't forget to count
↳ the terminating null)
8      echolen = strlen(argv[3]) + 1;
9      if (sendto(sock, argv[3], echolen, 0, (struct sockaddr
↳ *) &server_address, sizeof(server_address)) != echolen) {
10         perror("sendto");
11         exit(EXIT_FAILURE);
12     }
13
14     // Receive the message back from the server
15     if ((received = recvfrom(sock, buffer, 256, 0, (struct
↳ sockaddr *) &server_address, &addrlen)) != echolen) {
16         perror("recvfrom");
17         exit(EXIT_FAILURE);
18     }
19
20     cout << "Server (" << inet_ntoa(server_address.sin_addr)
↳ << ") echoed: " << received << " bytes: " << buffer << endl;
21
22     close(sock);
23     return 0;
24 }

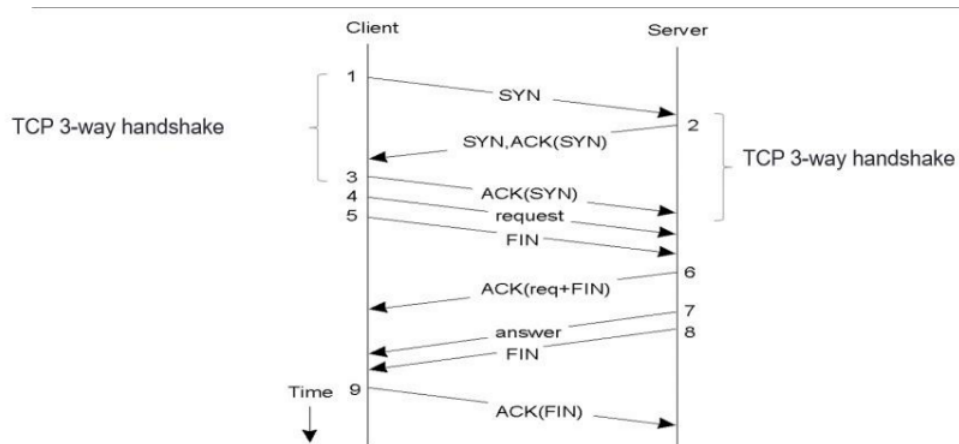
```

Transmission Control Protocol (TCP)

6.1 TCP / IP Protocol packet



6.2 TCP Communication



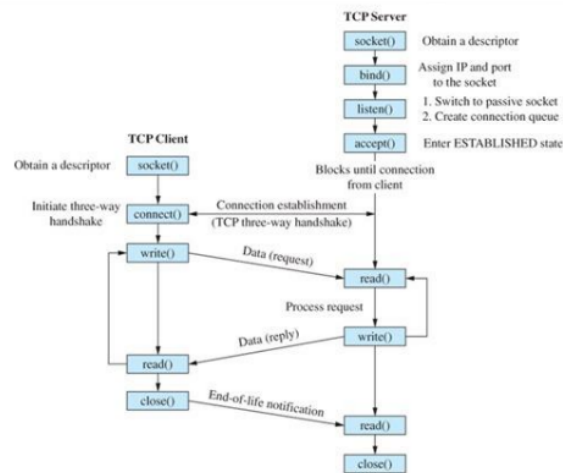
6.3 Socket System Calls

server

Primitive	Meaning
socket	Create a new communication endpoint
bind	Attach a local address to a socket
listen	Announce willingness to accept connections
accept	Block caller until a connection request arrives
connect	Actively attempt to establish a connection
write	Send(write) some data over the connection
read	Receive(read) some data over the connection
close	Release the connection

client

6.4 TCP Communications Pattern



6.5 System call: Socket

6.5.1 Signature

```
1 int socket(int domain, int type, int protocol)
```

- Creates a new socket, as end point to a communications link
- **Domain** is set to **AF_INET**

- **Type** is set to **SOCK_STREAM** for stream communications
- **Protocol** is set to 0, i.e 0 default TCP
- Returns socket descriptor
 - used in bind, listen, accept, connect, write, read, close

6.6 Client System call: connect

6.6.1 Signature

```
1  int connect(int sockfd, const struct sockaddr* addr, socklen_t  
    ↪  addrlen)
```

- connects socket to remote IP number and port
- **struct sockaddr** holds address information
 - will accept **struct sockaddr_in** pointer
- **addrlen** specifies length of **addr** structure
- returns 0 on success, -1 otherwise

6.7 TCP Client illustration

```

1  /*
2   * echoTCPClient.cxx
3   *
4   * TCP echo client
5   *
6   *      sends message to echo server
7   *      waits for message received from server
8   *
9   *      command line arguments:
10  *          argv[1] IP number of server
11  *          argv[2] port number to send to
12  *          argv[3] message to send
13  *
14  */
15 #include <sys/types.h>
16 #include <sys/socket.h>
17 #include <arpa/inet.h>
18 #include <unistd.h>
19 #include <netinet/in.h>
20
21 #include <cstdio>
22 #include <cstdlib>
23 #include <cstring>
24 #include <iostream>
25 using namespace std;
26
27 int main(int argc, char *argv[]) {
28
29     if (argc != 4) {
30         cerr << "USAGE: echoTCPClient server_ip port
↵ message\n";
31         exit(EXIT_FAILURE);
32     }
33
34     char buffer[256];
35     int echolen, received = 0;
36
37     int sock;
38     struct sockaddr_in server_address; // structure for
↵ address of server
39
40     // Create the TCP socket
41     if ((sock = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
42         perror("socket");
43         exit(EXIT_FAILURE);
44     }
45

```

```

1      // Construct the server sockaddr_in structure
2      memset(&server_address, 0, sizeof(server_address));
↳ /* Clear struct */
3      server_address.sin_family = AF_INET;
↳ /* Internet/IP */
4      server_address.sin_addr.s_addr = inet_addr(argv[1]);
↳ /* IP address */
5      server_address.sin_port = htons(atoi(argv[2]));
↳ /* server port */
6
7      // connect to server
8      if (connect(sock, (struct sockaddr *) &server_address,
↳ sizeof(server_address)) < 0) {
9          perror("connect");
10         exit(EXIT_FAILURE);
11     }
12
13     // Send the message to the server (don't forget to count
↳ the terminating null)
14     echolen = strlen(argv[3]) + 1;
15     if (write(sock, argv[3], echolen) < 0) {
16         perror("write");
17         exit(EXIT_FAILURE);
18     }
19
20     // Receive the message back from the server
21     if ((received = read(sock, buffer, 256)) < 0) {
22         perror("read");
23         exit(EXIT_FAILURE);
24     }
25
26     cout << "Server (" << inet_ntoa(server_address.sin_addr)
↳ << ") echoed: " << received << " bytes: " << buffer << endl;
27
28     close(sock);
29     return 0;
30 }

```


6.8 Server system call: bind

6.8.1 Signature

```
1  int bind(int sockfd, const struct sockaddr* addr, socklen_t  
    ↪      addrlen)
```

- assigns address to socket: IP number and port
- **struct sockaddr** holds address information
 - will accept **struct sockaddr_in** pointer
- **addrlen** specifies length of **addr** structure
- returns 0 on success, -1 otherwise

6.9 Server system call: Listen

6.9.1 Signature

```
1  int listen(int sockfd, int backlog)
```

- marks socket as passive socket
 - it will be used to accept incoming requests via accept
 - Term: "server socket"
- **backlog** specifies length of incoming connection queue
- returns 0 on success, -1 otherwise

6.10 Server system call: Accept

6.10.1 Signature

```
1  int accept(int sockfd, struct sockaddr* addr, socklen_t* addrlen)
```

- extracts connection request from incoming queue
- creates a new connected socket
 - returns a new file descriptor for that socket, returns -1 on failure
- **struct sockaddr** holds address information
 - will accept struct **sockaddr_in** pointer
- **addrlen** specifies length of **addr** structure

6.11 TCP Server illustration

```

1  /*
2   * echoTCPServer.cxx
3   *
4   * TCP echo server
5   *
6   *      loops/waits for message received from client
7   *      send message back to client
8   *
9   *      command line arguments:
10  *      argv[1] port number to receive from
11  *
12  */
13  #include <sys/types.h>
14  #include <sys/socket.h>
15  #include <arpa/inet.h>
16  #include <unistd.h>
17  #include <netinet/in.h>
18
19  #include <cstdio>
20  #include <cstdlib>
21  #include <cstring>
22  #include <iostream>
23  using namespace std;
24
25  int main(int argc, char *argv[]) {
26      // check arguments
27      if (argc != 2) {
28          cerr << "USAGE: echoTCPServer port\n";
29          exit(EXIT_FAILURE);
30      }
31
32      // Create the TCP socket
33      int sock = socket(AF_INET, SOCK_STREAM, 0);
34      if (sock < 0) {
35          perror("socket");
36          exit(EXIT_FAILURE);
37      }
38      // create address structures
39      struct sockaddr_in server_address; // structure for
↵ address of server
40      struct sockaddr_in client_address; // structure for
↵ address of client
41      unsigned int addrlen = sizeof(client_address);

```

```

1      // Construct the server sockaddr_in structure
2      memset(&server_address, 0, sizeof(server_address)); /*
↳ Clear struct */
3      server_address.sin_family = AF_INET; /*
↳ Internet/IP */
4      server_address.sin_addr.s_addr = INADDR_ANY; /*
↳ Any IP address */
5      server_address.sin_port = htons(atoi(argv[1])); /*
↳ server port */
6
7      // Bind the socket
8      if (bind(sock, (struct sockaddr *) &server_address,
↳ sizeof(server_address)) < 0) {
9          perror("bind");
10         exit(EXIT_FAILURE);
11     }
12
13     // listen: make socket passive and set length of queue
14     if (listen(sock, 64) < 0) {
15         perror("listen");
16         exit(EXIT_FAILURE);
17     }
18
19     cout << "echoServer listening on port: " << argv[1] <<
↳ endl;
20
21     // Run until cancelled
22     while (true) {
23         int connSock=accept(sock, (struct sockaddr *)
↳ &client_address, &addrlen);
24         if (connSock < 0) {
25             perror("accept");
26             exit(EXIT_FAILURE);
27         }
28         // read a message from the client
29         char buffer[1024];
30         int received = read(connSock, buffer,
↳ sizeof(buffer));
31         if (received < 0) {
32             perror("read");
33             exit(EXIT_FAILURE);
34         }
35         cout << "Client (" <<
↳ inet_ntoa(client_address.sin_addr) << ") sent " << received
↳ << " bytes: " << buffer << endl;
36         // write the message back to client
37         if (write(connSock, buffer, received) < 0) {
38             perror("write");
39             exit(EXIT_FAILURE);
40         }
41         close(connSock);
42     }
43
44     close(sock);
45     return 0;
46 }

```

TCP Server & Shell Job Control

7.1 Improve TCP Client

- Become Useful as generic client to any TCP server
- **Improvements:**
 - Accept FQDN as serve address
 - read and process complete server response

7.1.1 Accept FQDN as server address

```
1  // Lookup FQDN
2  struct addrinfo* res, hints;
3  memset(&hints,0,sizeof(hints));
4  hints.ai_family = AF_INET;
5  hints.ai_socktype = SOCK_STREAM;
6  int error = getaddrinfo(argv[1], argv[2], &hints, &res);
7
8  if (error) { ... }
9
10 // Create the tcp socket as normal
11
12 // Connect to server
13 if (connect(sock, res->ai_addr, res->ai_addrlen) < 0 ) { ... }
```

7.1.2 Process complete server response

```
1  // Recieve the message back from the server
2  do {
3      recieved = read(sock, buf, sizeof(buf))
4      if (recieved < 0) { ... }
5      cout.write(buf, recieved);
6  } while (recieved > 0);
```

7.1.3 Full example

```
1  /*
2  * TCPClient.cxx
3  *
4  * TCP client
5  *
6  *         sends message to TCP server
7  *         waits for message received from server
8  *
9  *         command line arguments:
10 *             argv[1] FQDN of server
11 *             argv[2] port number to send to
12 *             argv[3] request to send
13 *
14 */
15 #include <sys/socket.h>
16 #include <netdb.h>
17 #include <unistd.h>
18 #include <cstdlib>
19 #include <cstring>
20 #include <cstdio>
21 #include <iostream>
22 using namespace std;
23
24 int main(int argc, char *argv[]) {
25     if (argc != 4) {
26         cerr << "USAGE: TCPClient server_name port
↵ request\n";
27         exit(EXIT_FAILURE);
28     }
29     // lookup FQDN
30     struct addrinfo *res, hints;
31
32     memset(&hints, 0, sizeof(hints));
33     hints.ai_family = AF_INET;
34     hints.ai_socktype = SOCK_STREAM;
35
36     int error = getaddrinfo(argv[1], argv[2], &hints, &res);
37     if (error) {
38         cerr << argv[1] << ": " << gai_strerror(error)
↵ << endl;
39         exit(EXIT_FAILURE);
40     }
41     char buffer[1024];
42     int sent, received;
43 }
```

```

1      // Create the TCP socket
2      int sock = socket(AF_INET, SOCK_STREAM, 0);
3      if (sock < 0) {
4          perror("Failed to create socket");
5          exit(EXIT_FAILURE);
6      }
7
8      // connect to server
9      if (connect(sock, res->ai_addr, res->ai_addrlen) < 0) {
10         perror("connect");
11         exit(EXIT_FAILURE);
12     }
13
14     // Send the message string to the server
15     sent = write(sock, argv[3], strlen(argv[3])+1);
16     if (sent < 0) {
17         perror("write");
18         exit(EXIT_FAILURE);
19     }
20
21     // Receive the message back from the server
22     do {
23         received = read(sock, buffer, sizeof(buffer));
24         if (received < 0) {
25             perror("read");
26             exit(EXIT_FAILURE);
27         }
28         cout.write(buffer, received);
29     } while (received > 0);
30     cout << endl;
31
32     close(sock);
33 }

```

7.2 Improve TCP server

7.2.1 TCP Server Fork

- server starts loop
 - blocks on accept for connection from client
 - after accept:
 - * accept returns dedicated connection socket
 - * server forks into parent and child process
- parent process
 - closes dedicated connection socket
 - continues to block for next accept
- child process
 - serves client request
 - communicates with client via dedicated connection socket

7.2.2 TCP Server Fork basic logic

```
1  while (true) {
2      connSock = accept(sock, ...);
3      if (fork()) { // Parent
4          close(connSock);
5      } else { // Child
6          // Process client's request via connSock
7          ...
8      }
9  }
```


7.2.3 TCP Server Fork example

```
1  /*
2  * TCPServerFork.cxx
3  *
4  * TCP echo server
5  *
6  *      loops/waits/forks for message received from client
7  *      send message back to client
8  *
9  *      command line arguments:
10 *      argv[1] port number to receive from
11 *
12 */
13 #include <sys/types.h>
14 #include <sys/socket.h>
15 #include <arpa/inet.h>
16 #include <unistd.h>
17 #include <netinet/in.h>
18
19 #include <cstdio>
20 #include <cstdlib>
21 #include <cstring>
22 #include <iostream>
23 using namespace std;
24
25 void processClientRequest( int connSock) {
26     int received;
27     char buffer[1024];
28
29     // read a message from the client
30     if ((received = read(connSock, buffer, sizeof(buffer)))
↪     <= 0) {
31         perror("read");
32         exit(EXIT_FAILURE);
33     }
34
35     cout << "Client sent " << received << " bytes: " <<
↪     buffer << endl;
36
37     // write the message back to client
38     if (write(connSock, buffer, received) < 0) {
39         perror("write");
40         exit(EXIT_FAILURE);
41     }
42     close(connSock);
43     exit(EXIT_SUCCESS);
44 }
```

```

1  int main(int argc, char *argv[]) {
2      if (argc != 2) {
3          cerr << "USAGE: TCPServerFork port\n";
4          exit(EXIT_FAILURE);
5      }
6
7      // Create the TCP socket
8      int sock = socket(AF_INET, SOCK_STREAM, 0);
9      if (sock < 0) {
10         perror("socket");
11         exit(EXIT_FAILURE);
12     }
13     // create address structures
14     struct sockaddr_in server_address; // structure for
↪ address of server
15     struct sockaddr_in client_address; // structure for
↪ address of client
16     unsigned int addrlen = sizeof(client_address);
17
18     // Construct the server sockaddr_in structure
19     memset(&server_address, 0, sizeof(server_address)); /*
↪ Clear struct */
20     server_address.sin_family = AF_INET; /*
↪ Internet/IP */
21     server_address.sin_addr.s_addr = INADDR_ANY; /*
↪ Any IP address */
22     server_address.sin_port = htons(atoi(argv[1])); /*
↪ server port */
23
24     // Bind the socket
25     if (bind(sock, (struct sockaddr *) &server_address,
↪ sizeof(server_address)) < 0) {
26         perror("bind");
27         exit(EXIT_FAILURE);
28     }
29
30     // listen: make socket passive and set length of queue
31     if (listen(sock, 64) < 0) {
32         perror("listen");
33         exit(EXIT_FAILURE);
34     }
35
36     cout << "TCPServer listening on port: " << argv[1] <<
↪ endl;

```

```

1      // Run until cancelled
2      while (true) {
3          int connSock=accept(sock, (struct sockaddr *)
↪      &client_address, &addrlen);
4          if (connSock < 0) {
5              perror("accept");
6              exit(EXIT_FAILURE);
7          }
8          // fork
9          if (fork()) {                // parent process
10             close(connSock);
11         } else {                    // child process
12             processClientRequest(connSock);
13         }
14     }
15     close(sock);
16     return 0;
17 }

```

7.3 Improved TCP Server and Client example: List directory

- After accept, server forks to service client request
 - Parent process will loop to next accept
- Child process serves client request
 - Read directory path name from client
 - Open directory
 - Read directory entries, send file names to client
 - End process

7.3.1 Program

```

1      /*
2      * TCPServerReadDir.cxx
3      *
4      * TCP server
5      *
6      *      loops/forks to serve request from client
7      *      opens directory, sends back lines of file names
8      *      to client
9      *
10     *      command line arguments:
11     *      argv[1] port number to receive requests on
12     */
13     #include <sys/types.h>
14     #include <sys/socket.h>
15     #include <netinet/in.h>
16     #include <errno.h>
17     #include <dirent.h>
18     #include <unistd.h>
19     #include <cstdio>
20     #include <cstdlib>
21     #include <cstring>
22     #include <iostream>
23     using namespace std;
24
25     void processClientRequest(int connSock) {
26         int received;
27         char path[1024], buffer[1024];
28
29         // read a message from the client
30         if ((received = read(connSock, path, sizeof(path))) < 0)
31     ↪ {
32             perror("receive");
33             exit(EXIT_FAILURE);
34         }
35
36         cout << "Client request: " << path << endl;
37
38         // open directory
39         DIR *dirp = opendir(path);
40         if (dirp == 0) {
41             // tell client that an error occurred
42             strcpy(buffer, path);
43             strcat(buffer, ": could not open directory\n");
44             if (write(connSock, buffer, strlen(buffer)) < 0)
45     ↪ {
46                 perror("write");
47                 exit(EXIT_FAILURE);
48             }
49             exit(EXIT_SUCCESS);
50         }
51
52         // read directory entries
53         struct dirent *dirEntry;
54         while ((dirEntry = readdir(dirp)) != NULL) {
55             strcpy(buffer, dirEntry->d_name);
56             strcat(buffer, "\n");
57             if (write(connSock, buffer, strlen(buffer)) < 0)
58     ↪ {
59                 perror("write");
60                 exit(EXIT_FAILURE);
61             }
62         }
63     }

```

```

1  int main(int argc, char *argv[]) {
2
3      if (argc != 2) {
4          cerr << "USAGE: TCPServerReadDir port\n";
5          exit(EXIT_FAILURE);
6      }
7
8      // Create the TCP socket
9      int sock = socket(AF_INET, SOCK_STREAM, 0);
10     if (sock < 0) {
11         perror("socket");
12         exit(EXIT_FAILURE);
13     }
14     // create address structures
15     struct sockaddr_in server_address; // structure for
↪ address of server
16     struct sockaddr_in client_address; // structure for
↪ address of client
17     unsigned int addrlen = sizeof(client_address);
18
19     // Construct the server sockaddr_in structure
20     memset(&server_address, 0, sizeof(server_address)); /*
↪ Clear struct */
21     server_address.sin_family = AF_INET; /*
↪ Internet/IP */
22     server_address.sin_addr.s_addr = INADDR_ANY; /*
↪ Any IP address */
23     server_address.sin_port = htons(atoi(argv[1])); /*
↪ server port */
24
25     // Bind the socket
26     if (bind(sock, (struct sockaddr *) &server_address,
↪ sizeof(server_address)) < 0) {
27         perror("bind");
28         exit(EXIT_FAILURE);
29     }
30
31     // listen: make socket passive and set length of queue
32     if (listen(sock, 64) < 0) {
33         perror("listen");
34         exit(EXIT_FAILURE);
35     }
36
37     cout << "TCPServerReadDir listening on port: " <<
↪ argv[1] << endl;
38
39     // Run until cancelled
40     while (true) {
41         int connSock=accept(sock, (struct sockaddr *)
↪ &client_address, &addrlen);
42         if (connSock < 0) {
43             perror("accept");
44             exit(EXIT_FAILURE);
45         }
46         // fork 46
47         if (fork()) { // parent process
48             close(connSock);
49         } else { // child process
50             processClientRequest(connSock);
51         }

```